

---

# Layer-wise Adaptive KV Cache Compression for Training-Free Pythia-70M Inference

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

Autoregressive language-model inference is limited by the key-value (KV) cache, whose memory and attention cost grow with context length. This report presents a training-free layer-wise KV cache compression method for Pythia-70M. Instead of applying one uniform cache budget to every transformer block, we keep 80% of the cache in shallow layers, 50% in middle layers, and 70% in deep layers, using a deterministic sink-plus-recent token selector. We evaluate PG-19 and WikiText with cached continuation perplexity, TTFT, TPOT, throughput, cache reduction, and approximate decode-attention FLOPs. The method reduces cache tokens and estimated attention FLOPs by 28.4%. CPU speed is mixed: the layer-wise schedule is slightly faster on PG-19 but slower on WikiText, while uniform baselines show that implementation overhead and text slice both matter. Code is available at <https://github.com/lilfry09/Training-free-LM-inference-acceleration-group/tree/layerwise-compression>.

## 1 Introduction

During autoregressive decoding, transformer language models cache keys and values from previous tokens so that later tokens do not recompute the full prefix. The cache is useful, but it scales linearly with sequence length and therefore becomes a memory and attention bottleneck. For a model with  $N$  layers,  $H$  attention heads, head dimension  $d$ , and cache length  $L$ , the KV storage is proportional to  $2NLHd$  and the per-token attention score/value matmuls are proportional to  $4NLHd$ .

The course project requires a training-free acceleration method on Pythia-70M (Biderman et al., 2023), evaluated on datasets such as PG-19 and WikiText. We focus on sequence-dimension KV compression: reducing how many historical positions remain in the cache. The central question is whether a simple layer-wise budget can reduce cache and compute while preserving enough quality.

Our contributions are: (1) a reusable layer-wise KV cache compressor for HuggingFace causal language models; (2) PG-19 and WikiText evaluations with PPL, TTFT, TPOT, throughput, cache reduction, and approximate FLOPs; and (3) an honest analysis of the quality-speed trade-off, including cases where lower cache size does not translate into lower CPU wall-clock time.

## 2 Method

**Layer-wise cache budgets.** Pythia-70M has 6 transformer layers. We divide them by depth: shallow layers 0–1 keep 80% of cached tokens, middle layers 2–3 keep 50%, and deep layers 4–5 keep 70%. The motivation is that shallow layers often preserve token-level features, middle layers can tolerate more compression, and deep layers should retain enough context for final prediction.

**Token selector.** For each layer, the compressor keeps a small prefix of sink tokens and fills the remaining budget with the most recent tokens. If layer  $l$  has cache length  $L_l$  and keep ratio  $\rho_l$ , the kept length is

$$K_l = \max(K_{\min}, \lceil \rho_l L_l \rceil).$$

In our experiments  $K_{\min} = 16$  and the prefix length is 4. The method uses no gradients, no fine-tuning, and no changes to model weights.

**Inference procedure.** The prompt is first run with a full KV cache. The resulting cache is compressed layer by layer, then greedy decoding continues with the compressed cache. We also evaluate two uniform baselines: Uniform-50% and Uniform-67%. The latter matches the approximate average keep ratio of the layer-wise schedule, so it tests whether the depth-dependent allocation itself is useful.

**Metrics.** Cached continuation PPL scores true continuation tokens after a 96-token prefix. Speed uses greedy generation of 32 new tokens and reports TTFT, TPOT, throughput, and speedup over the dense cache. We report cache reduction from the average cached tokens per layer. Approximate decode-attention FLOPs per generated token are computed as  $4NHLd$ , where  $\bar{L}$  is the average cached tokens per layer.

## 3 Experiments

**Setup.** All experiments use the local Pythia-70M checkpoint, CPU inference, `torch.float32`, 96 prompt tokens, 32 generated tokens, and 2 timing repeats. PG-19 uses one local test sample as allowed by the assignment. WikiText uses the local WikiText-103 test split. The exact commands are:

```
python eval_quick.py -dataset pg19 -output_json results_pg19.json
python eval_quick.py -dataset wikitext -output_json results_wikitext.json
```

**Table 1:** PG-19 quick evaluation with Pythia-70M on CPU.

| Method      | PPL   | TTFT   | TPOT   | Tok/s | Speedup | Cache red. |
|-------------|-------|--------|--------|-------|---------|------------|
| Dense       | 37.50 | 0.0780 | 0.0176 | 51.31 | 1.00    | 0.0%       |
| Uniform-50% | 38.46 | 0.0723 | 0.0157 | 57.17 | 1.11    | 43.0%      |
| Uniform-67% | 36.08 | 0.0697 | 0.0183 | 50.10 | 0.98    | 27.8%      |
| LayerWise   | 39.25 | 0.0715 | 0.0175 | 51.95 | 1.01    | 28.4%      |

**Table 2:** WikiText quick evaluation with Pythia-70M on CPU.

| Method      | PPL   | TTFT   | TPOT   | Tok/s | Speedup | Cache red. |
|-------------|-------|--------|--------|-------|---------|------------|
| Dense       | 80.36 | 0.0764 | 0.0170 | 52.92 | 1.00    | 0.0%       |
| Uniform-50% | 84.97 | 0.0838 | 0.0156 | 56.41 | 1.07    | 43.0%      |
| Uniform-67% | 84.52 | 0.0633 | 0.0161 | 56.91 | 1.08    | 27.8%      |
| LayerWise   | 82.30 | 0.0925 | 0.0206 | 43.73 | 0.83    | 28.4%      |

**Table 3:** Cache and approximate decode-attention FLOPs.

| Method      | Avg. cache/layer | MFLOPs/token | Cache red. | FLOPs red. |
|-------------|------------------|--------------|------------|------------|
| Dense       | 111.50           | 1.370        | 0.0%       | 0.0%       |
| Uniform-50% | 63.50            | 0.780        | 43.0%      | 43.0%      |
| Uniform-67% | 80.50            | 0.989        | 27.8%      | 27.8%      |
| LayerWise   | 79.83            | 0.981        | 28.4%      | 28.4%      |

67 **Ablation and error analysis.** The layer-wise schedule consistently reduces cache size  
68 and estimated attention FLOPs by 28.4%. On PG-19, it gives a small speedup with a  
69 moderate PPL increase. On WikiText, it preserves PPL better than the uniform schedules  
70 but is slower in this CPU run. This suggests that the current implementation is memory- and  
71 FLOPs-efficient but not always wall-clock faster, because Python loop overhead, cache object  
72 conversion, and tensor slicing can dominate at this small model size. The uniform baselines  
73 are important: Uniform-50% is faster but hurts WikiText PPL more, while Uniform-67%  
74 can be fast but has less stable quality. The layer-wise schedule is therefore best interpreted  
75 as a conservative quality-preserving compression policy rather than a universal speedup.  
76 This negative timing case is kept in the report to avoid overstating the method beyond  
77 what the CPU evidence supports.

78 **Group contribution statement.** Current submitted workload: Fu Ruoyu designed the  
79 layer-wise schedule, implemented the compressor and evaluation script, ran PG-19/WikiText  
80 experiments, wrote the analysis, README, and report. Workload: 100%.

## 81 4 Conclusion

82 We implemented a training-free layer-wise KV cache compression method for Pythia-70M.  
83 The method reduces cache and approximate decode-attention FLOPs by 28.4% on both PG-  
84 19 and WikiText. The wall-clock speed results are mixed, which is an important limitation  
85 rather than a failed claim: at Pythia-70M scale on CPU, compression overhead can offset  
86 theoretical savings. Future work should move compression into fused kernels and test longer  
87 contexts where KV-cache costs dominate more clearly.

## 88 References

- 89 Biderman, S., Schoelkopf, H., Anthony, Q., Bradley, H., O’Brien, K., Hallahan, E., and others.  
90 Pythia: A suite for analyzing large language models across training and scaling. *ICML*, 2023.
- 91 Merity, S., Xiong, C., Bradbury, J., and Socher, R. Pointer sentinel mixture models.  
92 arXiv:1609.07843, 2016.
- 93 Rae, J. W., Potapenko, A., Jayakumar, S. M., and Lillicrap, T. P. Compressive transformers for  
94 long-range sequence modelling. *ICLR*, 2020.